

Alzheimer's Disease Detection Using Multimodal Modality



By:

Durab Haider

46614

Muhammad Umar

46952

Supervised by:

Mr. Mubariz Rehman

Faculty of Computing

Riphaah International University, Islamabad

Spring 2026

A Dissertation Submitted To

Faculty of Computing,
Riphah International University, Islamabad

As a Partial Fulfillment of the Requirement for the Award
of the Degree of
Bachelors of Science in Computer Science

Faculty of Computing
Riphah International University, Islamabad

Date: 28/04/2026

Final Approval

This is to certify that we have read the report submitted by ***Durab Haider (46614), Muhammad Umar (46952)***, for the partial fulfillment of the requirements for the degree of the Bachelors of Science in Computer Science (BSCS). It is our judgment that this report is of sufficient standard to warrant its acceptance by Riphah International University, Islamabad for the degree of Bachelors of Science in Computer Science (BSCS).

Committee:

1

Mr. Mubariz Rehman
(Supervisor)

2

Mr. Musharraf Ahmed
(Head of Department)

Declaration

We hereby declare that this document “**Alzheimer Disease Detection Using Multimodal Modality**” in the name of God, neither as a whole nor as a part has been copied out from any source. It is further declared that we have done this project with the accompanied report entirely on the basis of our personal efforts, under the proficient guidance of our teachers, especially our supervisor **Mr. Mubariz Rehman**. If any part of the system is proved to be copied out from any source or found to be reproduction of any project from anywhere else, we shall stand by the consequences.

Durab Haider

46614

Muhammad Umer

46952

Dedication

This project is dedicated to our parents for their unconditional love and support. We sincerely thank our supervisor “Mr. Mubariz Rehman” for his guidance and encouragement throughout our journey. Special appreciation goes to our friends for their motivation to help us keep going. This would not have been possible without the support of all who contributed directly or indirectly.

Acknowledgement

First of all, we are obliged to Allah Almighty the Merciful, the Beneficent and the source of all Knowledge, for granting us the courage and knowledge to complete this Project.

We would like to express our sincere gratitude to our supervisor “Mr. Mubariz Rehman” for their valuable guidance, continuous support, and constructive feedback throughout the course of Final Year Project. Their expertise and support were important to the successful completion of this work. We are really grateful to our parents who motivated us when times were challenging. We also respect the help and support of friends when required.

Durab Haider

46614

Muhammad Umar

46952

Abstract

Alzheimer's disease is an important illness affecting the brain, impacting memory, thinking and living day to day is common around the world and activities. Traditional methods for the diagnosis of Alzheimer's disease can take a long time, and is largely dependent on a doctor's expertise, and which may sometimes delay early detection, in this project, AI is being applied to aid in the speedy and correct diagnosis of Alzheimer's disease. The system analyzes brain MRI images along with cognitive tests to classify patients into different Cognitive Impairment, and Alzheimer's disease, are classified as stages of the disease. The stages of the disease are Cognitively Normal, Early Mild Cognitive Impairment, and Late Mild Cognitive Impairment, and Alzheimer's disease. This helps minimize the human impact. It brings error to light, and enables early diagnosis; this is vital to better treatment and care. Overall, this project aims to provide a reliable and easy-to-use tool that helps in early detection of Alzheimer and an ease for its users

Table of Contents

Chapter 1: Introduction	1
1.1 Goals and Objectives	4
1.2 Scope of the Project	4
Chapter 2: Literature Review	6
2.1 Introduction	7
2.2 Background and Problem Elaboration	7
2.3 Definition	8
2.4 Detailed Literature Review	8
2.4.1 Related Research Work 1: MRI-Based Deep Learning Approaches	9
2.4.2 Related Research Work 2: Multimodal Deep Learning Frameworks	9
2.4.3 Multimodal Interpretability and Sophisticated Multimodal Models	10
2.5 Literature Review Summary Table	10
2.6 Research Gap	12
2.7 Problem Statement	12
Chapter 3: Requirements and Design	14
3.1 Requirements	15
3.1.1 Functional Requirements	15
3.1.2 Non-Functional Requirements	16
3.1.3 Hardware and Software Requirements	16
3.2 Proposed Methodology	16
3.2.1 Overview	16
3.2.2 Steps	16
3.3 System Architecture	17
3.4 Use Cases	18
3.4.1 User	18
3.4.2 Admin	19
3.4.3 Complete Use Case Diagram	19
3.4.4 Fully Dressed Use Case	20
Chapter 4: Implementation and Test Cases	28
4.1 Implementation	29
4.1.1 Development Stack and Technical Environment	29
4.1.2 Imaging and Preprocessing Pipeline	29

4.1.3 Cognitive Vectorization and Fusion Logic	30
4.2 Test case Design and description	30
4.2.1 Test Case 1: User Registration and Authentication	31
4.2.2 Test Case 2: Multimodal Data Upload.....	31
4.2.3 Test Case 3: Result Display	32
4.2.4 Test Case 4: Admin User Management	33
4.3 Test Metric	35
4.3.1 Test Case Matric	35
Chapter 5: Experimental Results and Analysis.....	36
5.1 Introduction.....	37
5.2 Experimental Setup	37
5.3 Performance Result and Analysis	37
5.3.1 VGG16_ LeakyReLU	38
5.3.2 ResNet50_LeakyReLU	38
5.3.3 IncV3_MobNet_LeakyReLU	39
5.3.4 IncV3_DenseNet121.....	40
5.3.5 ResNet50_Cog	41
5.3.6 IncV3_MobNet_Attention_Pooling.....	42
5.3.7 IncV3_MobNet_Cog_Fusion_Softmax (Proposed Model)	43
5.3.8 Model Comparison.....	44
5.3.9 Detailed Analysis of Final Model	45
5.4 Analysis and Discussion	45
5.5 Conclusion	46
Chapter 6: Conclusion and Future Directions.....	47
6.1 Conclusion	48
6.2 Scope and Objective Assessment.....	48
6.3 Challenges Faced	49
6.4 Limitations	49
6.5 Future Work.....	49
References	51
Appendix.....	53
Appendix A: Dataset Samples	53
Appendix B: Multimodal Model Code Snippet	54

List of Tables

Table 2-1 -1 Summary Table (Literature Review)	11
Table 3-1 Functional Requirements of User	15
Table 3-2 Functional Requirements of Admin	15
Table 3-3 Non-Functional Requirements.....	16
Table 3-4 Use Case Table (User).....	18
Table 3-5 Use Case Table (Admin)	19
Table 3-6 User Login Fully Dressed Use Case Table.....	21
Table 3-7 Upload MRI Fully Dressed Use Case Table	21
Table 3-8 Cognitive Test Fully Dressed Use Case Table	22
Table 3-9 Diagnostic Results Fully Dressed Use Case Table.....	23
Table 3-10 Admin Login Fully Dressed Use Case Table	24
Table 3-11 View Users Fully Dressed Use Case Table.....	24
Table 3-12 Manage Users Fully Dressed Use Case Table.....	25
Table 3-13 Monitor MRI and Cog Data Fully Dressed Use Case Table	26
Table 3-14 Update Model Fully Dressed Use Case Table.....	27
Table 4-1 Test Case Table (User Registration).....	31
Table 4-2 Test Case Table (Multimodal Data Upload)	32
Table 4-3 Test Case Table (Result Display)	33
Table 4-4 Test Case Table (Admin Management).....	34
Table 4-5 Test Matric.....	35
Table 5-1 Confusion Matrix VGG16_LeakyReLU	38
Table 5-2 Evaluation Matric VGG16_LeakyReLU.....	38
Table 5-3 Confusion Matrix ResNet50_LeakyReLU	39
Table 5-4 Evaluation Matric ResNet50_LeakyReLU.....	39
Table 5-5 Confusion Matrix IncV3_MobNet_LeakyReLU	40
Table 5-6 Evaluation Matric IncV3_MobNet_LeakyReLU	40
Table 5-7 Confusion Matrix IncV3_DenseNet121.....	41
Table 5-8 Evaluation Matric IncV3_DenseNet121	41
Table 5-9 Confusion Matrix ResNet50_Cog	42
Table 5-10 Evaluation Matric ResNet50_Cog.....	42
Table 5-11 Confusion Matrix IncV3_MobNet_Att_Pooling.....	43
Table 5-12 Evaluation Matric IncV3_MobNet_Att_Pooling	43

Table 5-13 Confusion Matrix IncV3_MobNet_Cog_Fusion_Softmax	44
Table 5-14 Evaluation Matric IncV3_MobNet_Cog_Fusion_Softmax.....	44
Table 5-15 Model Comparison Table	45

List of Figures

Figure 3-1 Layered Architecture Diagram.....	17
Figure 3-2 User Use Case Diagram	18
Figure 3-3 Admin Use Case Diagram.....	19
Figure 3-4 Complete Use Case Diagram	20

Abstract

The proposed Alzheimer's disease Detection system will be intelligent and an efficient platform to help to identify and classify the onset of Alzheimer's in the early stages disease. The system is designed based on the latest deep learning technology to analyze brain MRI images along with cognitive assessment data for accurate diagnosis. By the model can be trained to classify images using image processing and data analysis techniques, with the help of AI. Patients into different stages such as Cognitively Normal, Mild Cognitive Impairment, and Alzheimer's disease. This automatic method decreases reliance on manual methods and makes the diagnostic process quicker and easier, with the help of minimizing human error more reliable.

The system is user friendly so that it is easy to interact with medical practitioners and scientists. It supplies simple and systematic results which are assist in making decisions in clinical settings. The AI-driven framework provides for high accuracy and scalability, appropriate for real world health care use settings. This project will help to identify disease at an early stage and improve disease stage, allowing the crop to be better managed. Facilitate prompt medical interventions, enhance the care of patients and participate in contributions to the creation of smart and eco-friendly healthcare solution.

Chapter 1: Introduction

Chapter 1: Introduction

Alzheimer disease (AD) is a progressive neurodegenerative condition and is the major cause of global dementia, and thus, a substantial factor of age-related cognitive impairment and functional loss. The disease is characterized by the progressive impairment of the memory, executive functions as well as behavioral bodies which eventually leads to loss of independence and high tendency of death. As the number of the aging population globally grows at an accelerated rate, AD is expected to increase significantly, burdening various health care systems and the relatives as well. Irrespective of intensive research work, effective disease-modifying measures are few making early and accurate diagnosis of the disease to be one of the key goals in managing Alzheimer disease.

The traditional form of diagnostic procedures mostly consists of clinical analysis, mental screening tests, and the visual analysis of neuroimaging examinations. Even though assessments of cognition, like the Montreal Cognitive Assessment (MoCA) and Clinical Dementia Rating Scale (CDR) are commonly used, they are inherently biased and prone to educational, cultural, and examiner bias significantly. Similarly, as a technique, the manual analysis of magnetic resonance images (MRI) is associated with labor-intensive processing and inter-observer variation hindering the reliability of this technique in the detection of subtle neurodegenerative markers of early-stage Alzheimer disease and mild cognitive impairment (MCI). As a result, there are a great number of people being diagnosed with one only when there is an irreversible neural damage, which can be very large.

The new developments in artificial intelligence (AI) and especially deep learning have shown immense potential to address all these diagnostic limitations. Convolutional neural networks (CNNs) have been demonstrated to be very effective in detecting discriminative features over high-dimensional medical images such as structural MRI scans and thus has been developed to detect disease-related patterns that traditional methods of analysis could not capture. However, one of the main drawbacks of many existing AI-based approaches is that they use single-modality data, and most often, it is MRI. The nature of the disease is that Alzheimer disease is fundamentally multifactorial and that structural brain changes and cognitive decrease interact in complex ways, therefore, the use of unimodal approaches may not be able to be considered as capturing the entire spectrum of the disease progression.

Multimodal theories of learning have now risen to importance in the quest to solve this issue through the incorporation of heterogeneous sources of data, including neuroimaging and cognitive measures, into a single utilized predictive model. Previous studies have confirmed the fact that a combination of MRI characteristics with clinical and neuropsychological data could significantly increase the accuracy and strength of classification of individuals with cognitive normalcy, early MCI, late MCI, and patients with Alzheimer's disease. However, multimodal fusion is associated with various technical issues such as the heterogeneity of data, data alignment, class imbalance, and explainability of models. Most of the existing researches use simplistic fusion models or do not thoroughly compare their outcomes to all clinically relevant disease stages and thus are unable to be applied in translation.

This project aims to fill such gaps by suggesting a multimodal deep learning model to detect early signs of the Alzheimer disease with the integration of the MRI-based features and cognitive tests scores into a single attention-based model. Deep CNNs are used to process structural MRI volumes to obtain high levels of spatial representations, whereas cognitive measures of MoCA, ADAS-13, and CDR-SB are used as additional clinical measures. The attention mechanisms are used to enhance feature learning by highlighting the diagnostically relevant areas and patterns to enhance the performance of classification and generalization.

The major advantage of this project is its end-to-end system design, which is not contained in model development, but also in a realistic user interface (UI). The enacted UI allows uploading MRI data and cognitive tests thus enabling real time prediction and visualization of progressions of diagnosis. This type of system level integration is the key to closing the gap between the academic research and clinical implementation, but it has not been studied in detail in many extant studies. The proposed solution would lead to better usability and possible adoption in the real-world healthcare environment by combining the forefront usability with the backend intelligence.

Besides, the project uses a rigorous appraisal model based on practical neuroimaging datasets that one can obtain via novel repositories like the Alzheimer Disease Neuroimaging Initiative (ADNI). To provide clinical relevance and strength, performance is measured using several diagnostic categories where one of the categories is Alzheimer disease, the other is early MCI, the third is late MCI and fourth who are cognitively normal. Although this study accepts limitations to diversification

between datasets and demographic biases, reproducibility, scalability and methodological transparency are of the highest priority.

Conclusively, this study is relevant to the field of AI-based medicine because it provides a powerful, multimodal, and explainable deep learning model of Alzheimer disease identification. The combination of MRI and cognitive data, exploitation of attention-based architectures, and deployment of a user interface make the proposed framework contribute to the future development of diagnostic methodologies and provide the basis of a future clinical decision-support system.

1.1 Goals and Objectives

The following goals are aimed at with this project:

1. Developing and establishing a multimodal Deep Learning program, combining structural MRI data and cognitive assessment data to detect Alzheimer disease.
2. To test and evaluate attention boosted CNN architectures having ability to identify discriminative neuroimaging markers that can be used in disease progression.
3. To classify the subjects into clinically meaningful groups, e.g., AD, LMCI, EMCI, and CN.
4. To develop a usability interface which will allow user-friendly interacting with the diagnostic system, supporting usability and real-time prediction in the system.
5. To evaluate the performance of the system critically using pre-established databases.

1.2 Scope of the Project

This project focuses on the design, development and testing of a multimodal deep learning system for early identification and classification of Alzheimer's disease using MRI and cognitive assessment data. The project provides a solution in the diagnostic problem by applying an integrated method based on computational methods consisting of medical imaging, artificial intelligence, clinical data and clinical studies.

The project will consist of preprocessing and analyzing the brain MRI scans, obtained out of the private datasets, to determine patterns of structures related to the cases of Alzheimer's disease and its preliminaries. Deep convolutional neural networks are used

to learn imaging discriminative features automatically without human intervention. Attention mechanisms are integrated to increase the feature significance through emphasizing the areas of the brain that are diagnostically important, hence making the classification more robust and interpretable.

Simultaneously, the scope encompasses the incorporation of common measures of cognitive assessment, including MoCA, ADAS-13, and CDR-SB, which can be viewed as complementary items of cognitive deterioration. Those cognitive characteristics are combined with MRI-based representations in a multimodal framework of learning to enhance predictive performances in clinically-significant subsets, such as cognitively normal (CN), early mild cognitive impairment (EMCI), late mild impaired (LMCI), and Alzheimer disease (AD).

The project also involves deployment of various deep learning architectures to make a comparative analysis of the various architectures, like CNN based MRI models and multimodal fusion networks and measure the performance in terms of established metrics like accuracy, precision, recall and F1-score. The model training, validation and testing is carried out through controlled experimental setting so as to achieve reproducibility and analytical rigor.

Furthermore, development of an active user interface where users will be able to post MRI images and cognitive inputs and get automatic diagnostic forecasts is also within the scope. This element connects the research gap between developing an algorithm and usability, which helps to consider potential clinical or research-based implementation.

The present project is also confined to the diagnostic prediction and classification, and not the treatment recommendation, disease prognosis, or longitudinal progression modeling. Moreover, real-time clinical verification and hospital implementation are not within the bounds of the current research and are defined as the future work direction possibilities.

Chapter 2: Literature Review

Chapter 2: Literature Review

Here we will discuss about the literature review.

2.1 Introduction

Alzheimer disease (AD) is a progressive neurodegenerative disease and is one of the most urgent problems of the twenty-first century of the public-health sector. The pathology evolves in a stealthy manner and in most cases, it begins several years before the manifestation of clinical symptoms thus making the early diagnosis a daunting task. The traditional diagnostic paradigms have been based to a large extent on cognitive screening tools and the subjective interpretation of neuroimaging by clinicians. These biases based on subjective measurements often cause a failure to recognize even the smallest pathological changes at the earliest stages of the disease, which means that in the end it becomes the significant percentage of the patients who will only be diagnosed at the moment when it is already too late, and the neural damage cannot be removed. The recent advances in the field of artificial intelligence, specifically in deep learning (DL), have relegated medical image-related analysis processes and clinical decision-making workflows to a radically new stage. DL models have the ability to automatically derive patterns of discrimination in high-dimensional neuroimaging scans and thus perform better than conventional machine-learning methods that require hand-defined features [1]. Furthermore, recent studies have also been more specifically insightful to multimodal DL systems, in which magnetic resonance imaging (MRI) data are utilized together with cognitive and clinical evaluations to represent the characteristics of the disease complementarily. The chapter is a critical review of the modern literature on DL-based AD detection, specifically on multimodal frameworks, thus outlining their accomplishments, weaknesses, and outstanding gaps.

2.2 Background and Problem Elaboration

The pathological characteristics of AD cover several dimensions, including structural brain shrinkage, functional loss and cognitive impairment. Structural MRI is the most commonly used modality in the visualization of anatomical changes as hippocampal or cortical thinning, although it may be extraordinary fine at the early disease stage, and can have a significant inter-individual variability. Cognitive batteries (Montreal Cognitive Assessment, the Alzheimer Disease Assessment Scale-Cognitive subscale,

and the Clinical Dementia Rating-Sum of Boxes) offer independent information on functional deterioration, but are also bound to be affected by education levels at baseline, judge bias, and compliance of patients. Some of these limitations are alleviated in DL methodologies whereby, automated and data driven development of features can be performed using complex inputs. Architectures based on convolutional neural network (CNN) have shown high performance in diagnosis of AD using MRI data especially in the discrimination of AD patients and cognitively normal subjects [2]. However, fine-grained classification tasks, including distinguishing between early mild impairment of the condition, the mild (MCI) and late MCI in multimedia unimodal MRI-only systems often fail because of the similarity of structural responses. Therefore, constructing multimodal fusion methods that merge imaging, cognitive and clinical data is a method that is progressively applied by researchers and these methods promise to improve diagnostic sensitivity and strength [3]. Notwithstanding such developments, multimodal systems introduce new challenges such as variance of data, increased architectural complexity, inability to interpret, and reduced generalizability across cohorts. Such considerations spur an extensive and pertinent analysis of research that has been published.

2.3 Definition

Deep Learning (DL) is a branch of machine learning that uses multilayered neural networks that can learn hierarchical representations by directly operating on raw data. The major application of DL in Alzheimer research is in the form of CNNs, attention-based networks and imaging and clinical data analysis hybrid architectures [1]. Multimodal Learning refers to an approach to using heterogeneous data (i.e., MRI images and cognitive scores) to generate the same predictive framework. Multimodal learning aims at capitalizing on the complementary information that cannot be realized on a single modality [3]. Attention mechanisms are neural elements, which dynamically emphasize salient characteristics or modalities during training thus boosting model execution and providing restricted interpretability by illuminating diagnostically significant data [4].

2.4 Detailed Literature Review

Below is the detailed literature review.

2.4.1 Related Research Work 1: MRI-Based Deep Learning

Approaches

The literature that focused on AD detection using structural MRI has continued to grow after 2020 with the focus on DL models. Mohammed et al. (2024) review the state of the art of the DL based MRI classification models and find that CNN architectures significantly outperform the traditional classifiers in recognizing the disease related structures [1]. The authors, nevertheless, stress that the majority of MRI-only systems are based on large labelled datasets and are characterized by a reduced generalization when imposed on unseen cohorts. A CNN -based MRI classification model suggested by Sorour et al. (2024) demonstrated excellent results in the initial AD detection, thus, proving the effectiveness of automated feature extraction [2]. Although these results are encouraging, the study is not expanded following imaging data and lacks the input of cognitive or the clinical data to draw conclusions on the applicability of the findings to a real-life diagnostic situation. On the same note, Raj et al. (2025) proposed a hybrid CNN that includes the InceptionV3 and ResNet-50 in the classification of Alzheimer. Although the model achieved high accuracy, it was more difficult to display the middle diseases stage which highlights the inherent weakness of unimodal imaging techniques in any fine-grained classification [3]. Taken together, these studies demonstrate MRI-based DL models are potent structural pattern recognition tools with a limitation in the form of the inability to combine both functional and clinical aspects of the Alzheimer disease.

2.4.2 Related Research Work 2: Multimodal Deep Learning

Frameworks

In order to overcome the drawbacks of unimodal systems, recent studies are more likely to use multimodal DL methodologies. Golovanevsky et al. (2022) proposed an attention-based multipurpose model that combines MRI and clinical data, which has better multi-class classification accuracy and indicates the advantage of cross-modal interaction [4]. However, their method requires careful alignment of features and is overwhelmed by greater computing complexity. Amine (2025) performed a review of multimodal DL models to detect early-stage Alzheimer, and found that the combination of imaging and cognitive information is always improved as compared to single-modality methods in the diagnosis of Alzheimer in its domains of early susceptibility.

The review, however, points out that most multimodal systems are based on the wider use of preprocessing and feature engineering, which suggests issues of scalability and reproducibility. Wang et al. (2024) came up with the concept of a hybrid multimodal model involving the integration of MRI and cognitive tests to forecast the clinical development of dementia. Though the model has been effective in illustrating the relationship between structural changes and cognitive decline, it was evaluated using one dataset alone thus restraining conclusions on matters of generalization [5]. Their multimodal neural network incorporated MRI, PET and neuropsychological scores, creating a lightweight, efficient neural network with competitive results (Gupta et al., 2025). Nonetheless, the practical applicability is reduced based on dependence on PET imaging because of its high cost and unavailability in standard clinical practice [6]. According to these studies, the concept of multimodal learning enhances the accuracy of the diagnosis and staging of the disease.

2.4.3 Multimodal Interpretability and Sophisticated Multimodal Models

Recent sources emphasize the relevance of interpretability to the diagnostic systems based on the DL. One of the multimodal frameworks suggested by Castellano et al. (2024), which involves the use of both MRI and PET, has shown improved performance on detection of it but not regarding transparency of decisions made [7]. Raza et al. (2025) maintain that attention-based and explainable architectures are vital to clinical trust but cite that most of the existing models offer partial interpretability only [8]. Song et al. (2021) showed multimodal imaging fusion methods that retain modality-specific information and thus help to interpret the image but do not change performance [9]. Recent reviews incriminate transformer based multimodal architectures as a promise but an expensive and data intensive form of computation, which limits their utilization in resource constrained environments [10], [11].

2.5 Literature Review Summary Table

Below is the literature review table.

Table 2-1 -1 Summary Table (Literature Review)

No.	Study	Modalities	Method	Strengths	Limitations	Accuracies
1.	Mohammed et al. (2024)	MRI	DL Review	Comprehensive analysis	Unimodal focus	98.24%
2.	Sorour et al. (2024)	MRI	CNN	Strong early detection	No clinical data	99.92%
3.	Raj et al. (2025)	MRI	Hybrid CNN	High accuracy	Poor stage separation	96.6%
4.	Golovanevsky et al. (2022)	MRI + Clinical	Attention-based fusion	Improved multiclass	High complexity	96.88%
5.	Wang et al. (2024)	MRI + Cognitive	Hybrid DL	Predicts progression	Limited validation	90%
6.	Gupta et al. (2025)	MRI + PET + Cognitive	Lightweight DL	Efficient fusion	PET dependency	98.0%
7.	Castellano et al. (2024)	MRI + PET	CNN	Enhanced accuracy	Low interpretability	95.3%
8.	Raza et al. (2025)	Multimodal	Review	Highlights explainability	Conceptual	92.14%
9.	Song et al. (2021)	MRI + PET	Fusion DL	Preserves modality info	Limited staging	94.11%
10.	Mohsen (2025)	MRI	Review	Comparative insights	Dataset bias	97.5%
11.	Nagarajan et al. (2025)	MRI	Survey	Comprehensive scope	Lacks implementation	99.21%

2.6 Research Gap

The literature contains gaps. Even though there has been a lot of progress, some gaps still remain.

- Lack of interpretability of deep learning predictions to the clinical use.
- Weakness in extensive generalization of heterogeneous datasets.
- Algorithms on costly or unavailable modalities.
- Absence of end-to-end systems to integrate prediction and usable interfaces.

2.7 Problem Statement

Alzheimer's disease (AD) is a progressive neurodegenerative disorder and one of the greatest of three causes of dementia across the world. It has a major impact on patients, family members, and healthcare systems. Manual analysis of neuroimaging data makes early and accurate diagnosis of the mild cognitive impairment (MCI) stage difficult due to minor structural brain changes. Limited access to expert radiologists and advanced diagnostic tools is a challenge and not easily available in many health care facilities, especially in resource-limited settings, leading to delayed or inaccurate diagnosis.

Magnetic Resonance Imaging (MRI) is a widely used technique for the assessment of Alzheimer's disease. However, the current diagnostic methods have several major problems. The structural changes in the brain in AD differ in size, location and magnitude between patients, making it difficult to accurately determine the disease-relevant areas. There is also inter-subject variability, differences in scanning protocols and also noise in MRI data which have negative impact on automated analysis. These factors make it difficult to develop discriminative features for reliable and accurate disease classification.

Due to their improved performance in classifying Alzheimer's disease, deep neural networks (DNNs) have found impressive success in a number of applications; however, there is a considerable gap between the beginning of their first and second symptoms. Some of the disadvantages are the training times, the computational expense of training, and a heavy reliance on huge, well-trained datasets.

In real world, these requirements might not be feasible. Models may perform poorly in complex situations and are prone to overfitting with small data sets. Therefore, it is

difficult to deploy end-to-end deep learning models in real time or for large-scale applications. Clinical use still remains a major challenge.

There is a critical need for an efficient and reliable Alzheimer's disease diagnosis framework that can overcome these limitations by implementing robust deep feature extraction, reducing model training complexity, and improving generalization across subjects. This project aims to develop a deep learning-based system that utilizes pre-trained convolutional neural networks for MRI feature extraction, followed by optimized classification and multimodal fusion techniques. The proposed approach seeks to achieve accurate and early detection of Alzheimer's disease while minimizing computational overhead, thereby supporting clinicians with a practical, cost-effective, and decision-support tool for improved patient care.

Chapter 3: Requirements and Design

Chapter 3: Requirements and Design

This chapter aims at clearly defining the needed system functionality and its structural and functional design in terms of making it reproducible by the readers with a suitable technical background. The chapter establishes functional and non-functional requirements, hardware and software dependencies, proposed methodological workflow, system architecture and exemplary use cases. It focuses on modularity, interpretability, and clinical relevance, as it is proposed to be done in medical artificial intelligence system design best practice.

3.1 Requirements

Below is the list of functional and non-functional requirements of our system.

3.1.1 Functional Requirements

Here are the functional requirements of our system.

3.1.1.1 User

Here are the functional requirements of User.

Table 3-1 Functional Requirements of User

ID	Requirements
FR-1.1	User shall be able to create profile
FR-1.2	User shall be able to login
FR-1.3	User shall be able to upload MRI scan
FR-1.4	User shall be able to take cognitive test
FR-1.5	User shall be able to view diagnostic result on their dashboard

3.1.1.2 Admin

Here are the functional requirements of Admin.

Table 3-2 Functional Requirements of Admin

ID	Requirements
FR-2.1	Admin shall be able to login to system
FR-2.2	Admin shall view the list of registered users
FR-2.3	Admin shall be able to manage user accounts
FR-2.4	Admin shall be able to monitor MRI and cognitive data
FR-2.5	Admin shall be able to update model if needed

3.1.2 Non-Functional Requirements

Here are the non-functional requirements of our system.

Table 3-3 Non-Functional Requirements

Non-Functional Requirement	Description
Performance Efficiency	The system shall process individual MRI scans within an acceptable response time.
Usability	The system shall provide an intuitive user-friendly interface.

3.1.3 Hardware and Software Requirements

This system needs proper computational facilities to facilitate deep learning inference and processing data including:

- Multi-core CPU
- High end GPU
- 32GB RAM

These requirements are congruent to current research settings that are used in the development of AI in the medical profession.

3.2 Proposed Methodology

Here we will discuss about the proposed methodology.

3.2.1 Overview

Our system aims to detect and classify stages of Alzheimer's by using a web application integrated with deep learning model. Our system provides user friendly interface for users to verify the stage of Alzheimer effectively.

3.2.2 Steps

1. Data Collection: We collected the MRI dataset from ADNI which was not publicly available and was gotten upon request.
2. Data Preprocessing: Dataset was preprocessed from 3D to 2D with 35 center slices for each volume with dimensions 128x128. The dataset was split into train (70%), validation (20%), and test (10%).

3. **Model Training and Evaluation:** Train the CNN model using preprocessed dataset. The evaluation is done through confusion matrix.
4. **Integration:** Developing a web application which ensures user-friendly interface for users. Then we will integrate the model in our web app that will then do the detection based on the uploaded MRI and cognitive test taken and the system will generate the result on time.
5. **Testing and Validation:** Thorough testing will be done to ensure accuracy and reliability of our system and different CNN's will be trained to ensure we choose the best model.

3.3 System Architecture

A modular, layered approach is used to develop the system architecture to help it to be scalable and maintainable. The architecture is as follows:

1. **Data Input:** The user uploads the MRI and takes cognitive test, which will work as input to system for classification.
2. **Application Interface:** The web application provides a user-friendly interface that allows easy submission of MRI images and cognitive test. It communicates with the model and ensures smooth interaction for user.
3. **Deep Learning Model:** The system uses a hybrid multimodal model that analyzes MRI images and cognitive tests. It processes the inputs and classifies the patient into different categories such as CN, EMCI, LMCI, or AD.
4. **Results:** The classification results are presented to the user, showing the predicted stage.




Figure 3-1 Layered Architecture Diagram

3.4 Use Cases

Below are the details of the use cases, complete use case and fully dressed use case.

3.4.1 User

Below is the User use case.

Table 3-4 Use Case Table (User)

Actors	User
Pre-Conditions	User must have MRI images ready User must be registered in the system
Main Flow	User registers/logs-in to the system User uploads MRI brain scans User takes cognitive test User views diagnostic result on their dashboard
Post-Conditions	Uploaded data is stored securely for future reference




Figure 3-2 User Use Case Diagram

3.4.2 Admin

Below is the use case for Admin.

Table 3-5 Use Case Table (Admin)

Actors	Admin
Pre-Conditions	Admin must have login credentials.
Main Flow	Admin logs into the admin dashboard Admin views the list of registered users Admin can manage user accounts (update or delete account) Admin can monitor MRI and cognitive data Admin can update model if needed
Post-Conditions	Admin maintains system integrity




Figure 3-3 Admin Use Case Diagram

3.4.3 Complete Use Case Diagram

Below is the complete Use Case Diagram.




Figure 3-4 Complete Use Case Diagram

3.4.4 Fully Dressed Use Case

Below are the fully dressed use cases.

3.4.4.1 User Login

Below is the fully dressed use case for login.

Table 3-6 User Login Fully Dressed Use Case Table

Name	User Registration / Login		
Actor	User		
Summary	Allows a new user to create an account or an existing user to access their private dashboard.		
Pre-Conditions	User has a stable internet connection and access to the web portal.		
Post-Conditions	User is authenticated and redirected to their personal dashboard.		
Basic Flow			
Actor Action		System Response	
1	User enters credentials (Email/Password) and clicks Login.	2	System validates the credentials against the user database.
		3	System generates a secure session and grants access.
4	User views their personalized dashboard and past records.		

3.4.4.2 Upload MRI

Below is the fully dressed use case of MRI Upload.

Table 3-7 Upload MRI Fully Dressed Use Case Table

Name	Upload MRI Scans		
Actor	User		
Summary	Allows the user to upload MRI Scans.		
Pre-Conditions	User is logged in and has a valid MRI file.		
Post-Conditions	The file is successfully stored in the backend for processing.		
Basic Flow			
Actor Action		System Response	
1	User selects the MRI file and clicks the upload button.	2	System checks file format compatibility

		3	System saves the file and confirms upload success.
4	User confirms and proceeds to cognitive test.		

3.4.4.3 Take Cognitive Test

Below is the fully dressed use case of Cognitive Test.

Table 3-8 Cognitive Test Fully Dressed Use Case Table

Name	Take Cognitive Test		
Actor	User		
Summary	Allows the user to take cognitive test.		
Pre-Conditions	User has uploaded the MRI successfully.		
Post-Conditions	Cognitive test has been done and ready for multimodal fusion.		
Basic Flow			
Actor Action		System Response	
1	User answers the questions.	2	System validates that the scores are within the range.
		3	System maps the score to a 30-bit binary vector
4	User proceeds to final diagnostic request.		

3.4.4.4 View Diagnostic Result

Below is the fully dressed use case of View Diagnostic Result.

Table 3-9 Diagnostic Results Fully Dressed Use Case Table

Name	View Diagnostic Result		
Actor	User		
Summary	Allows the user to see the AI's classification based on the fusion of MRI and cognitive data.		
Pre-Conditions	Both MRI and Cognitive data have been submitted and processed.		
Post-Conditions	User receives a diagnostic forecast (e.g., AD, CN, EMCI, or LMCI).		
Basic Flow			
Actor Action		System Response	
1	User clicks the "Generate Result" button on the dashboard.	2	System runs the multimodal fusion model inference.
		3	System displays the diagnostic label.
4	User reads the detailed result on their dashboard.		

3.4.4.5 Admin Login

Below is the fully dressed use case of Admin Login.

Table 3-10 Admin Login Fully Dressed Use Case Table

Name	Admin Login		
Actor	Admin		
Summary	Provides the administrator with secure access to the system management tools.		
Pre-Conditions	Admin account exists with administrative privileges.		
Post-Conditions	Admin is redirected to the management panel.		
Basic Flow			
Actor Action		System Response	
1	Admin enters administrative credentials and clicks Login.	2	System verifies admin status and validates credentials.
		3	System provides access to user lists and model controls.
4	Admin views the system health and user overview.		

3.4.4.6 View Registered Users

Below is the fully dressed use case of View Registered Users.

Table 3-11 View Users Fully Dressed Use Case Table

Name	View Registered Users
Actor	Admin
Summary	Allows the admin to oversee all active accounts within the Alzheimer's Diagnostic System.
Pre-Conditions	Admin is successfully logged in to the admin dashboard.

Post-Conditions		Admin can see a comprehensive list of all users and their status.	
Basic Flow			
Actor Action		System Response	
1	Admin clicks on the "User Management" tab.	2	System fetches all user records from the primary database.
		3	System renders the list with search and filter options.
4	Admin identifies the users currently utilizing the platform.		

3.4.4.7 Manage User Accounts

Below is the fully dressed use case of Manage User Accounts.

Table 3-12 Manage Users Fully Dressed Use Case Table

Name		Manage User Accounts	
Actor		Admin	
Summary		Allows the admin to update or delete users.	
Pre-Conditions		Admin is viewing the list of registered users.	
Post-Conditions		User records are modified or permanently removed from the system.	
Basic Flow			
Actor Action		System Response	
1	Admin selects a user and clicks "Update" or "Delete".	2	System prompts for a second confirmation to prevent error.
		3	System executes the database query to modify/remove the record.
4	Admin receives a notification that the change is finalized.		

3.4.4.8 Monitor MRI and Cognitive Data

Below is the fully dressed use case of Monitor MRI and Cognitive Data.

Table 3-13 Monitor MRI and Cog Data Fully Dressed Use Case Table

Name		Monitor MRI and Cognitive Data	
Actor		Admin	
Summary		Allows the admin to monitor the type and volume of data being processed by the AI.	
Pre-Conditions		Admin is logged in and system data is being actively collected.	
Post-Conditions		Admin can ensure that data heterogeneity and imbalances are tracked.	
Basic Flow			
Actor Action		System Response	
1	Admin navigates to the "Data Analytics" section.	2	System displays statistics on uploaded scans and test scores.
		3	System highlights any anomalies or data processing errors.
4	Admin ensures the system is handling data partitions correctly.		

3.4.4.9 Update Model

Below is the fully dressed use case of Update Model.

Table 3-14 Update Model Fully Dressed Use Case Table

Name		Update Model	
Actor		Admin	
Summary		Allows the admin to replace the current AI model with a newly trained version to improve accuracy.	
Pre-Conditions		Admin has a new validated weights file (.h5) ready for deployment.	
Post-Conditions		The backend begins using the updated model for all new user requests.	
Basic Flow			
Actor Action		System Response	
1	Admin uploads the new model file in the settings panel.	2	System verifies the file structure and model architecture.
		3	System hot-swaps the model files in the backend directory.
4	Admin receives a success message indicating the new model is live.		

Chapter 4: Implementation and Test Cases

Chapter 4: Implementation and Test Cases

In this chapter, the detailed implementation decisions that support the Alzheimer disease (AD) detection system are identified, and the architecture of the prototype that enables other professionals to reproduce the pipeline and results thereof are explained.

4.1 Implementation

This section details how the multimodal framework was converted into a functional software system. The implementation is based on the smooth integration between the deep learning backend and a modern web front-end.

4.1.1 Development Stack and Technical Environment

To ensure the system is both powerful and accessible, we utilized a decoupled architecture:

1. Backend & AI: Node.js was used for backend, model was developed in Python using TensorFlow and Keras for execution. We used FastAPI to handle requests between the AI and the web interface.
2. Frontend Dashboard: React.js was used for frontend development providing a highly responsive and intuitive user interface.

4.1.2 Imaging and Preprocessing Pipeline

The imaging pipeline ensures that raw brain scans are compatible with the neural network's input shape.

1. Data Handling: The system communicates with 3D NIFTI volumes via the Nibabel library. It detects and extracts the main axial slices, which contain the most important indicators of hippocampal shrinkage.
2. Normalization: We implemented Z-score Intensity Normalization and resized the slices to a uniform 128x128 pixels. This will ensure that differences in MRI are kept to a minimum scanner hardware and do not affect the model's accuracy.
3. Hybrid Feature Extraction: MRI branch performs a dual-stream feature extraction ensemble. InceptionV3 captures multi-scale spatial details and MobileNet offers efficient and high-level structural patterns that ensure the model is easy to deploy in the web with a lightweight footprint.

4.1.3 Cognitive Vectorization and Fusion Logic

To incorporate clinical data, I developed a custom feature-engineering pipeline for the cognitive assessments:

1. **30-Question Vector:** The model does not get raw scores from MoCA, ADAS-13, and CDR-SB. Thus, I created a logic-gate function that converts these scores into a binary vector with 30 bits. Memory, orientation, and executive function are three separate cognitive domains that are represented as active or inactive bits.
2. **Attention-Based Fusion:** When combined with backbone attentive fusion, the model can focus on specific areas of the image instead of the entire structure. Attention weights are determined by the system for that specific patient once the MRI features and cognitive vector converge at the fusion head to prioritize specific data, leading to the high accuracy for the final testing phase.
3. **Web Integration and Interface Workflow:** The React dashboard design guarantees that non-technical users can access the complex backend.
4. **Asynchronous Prediction:** The React frontend sends an asynchronous request when the user uploads an MRI file and inputs cognitive scores.
5. **Inference Management:** The trained weights are loaded in the backend (incv3_mobnet_cog30_fusion_model.h5). It makes the prediction across the four classes (AD, CN, EMCI, and LMCI).
6. **Result Delivery:** The prediction is given back to the dashboard within 1 to 5 seconds. The interface then displays the diagnostic forecast alongside the classification label, completing the end-to-end diagnostic support loop.

4.2 Test case Design and description

Here, we have given some test cases to check if the Alzheimer's Detection system is working fine or not. These tests were designed to see how well the system coped with brain scans and cognitive test scores at the same time. We have tested all the main points, from uploading the data and processing the images to the final prediction and how the user interface looks. We used standard testing methods to make sure that the system is not only accurate in theory but also stable and easy to use for anyone trying it.

4.2.1 Test Case 1: User Registration and Authentication

Below is the test case of User Registration and Authentication.

Table 4-1 Test Case Table (User Registration)

Authentication Module			
UC-01 User Login/Registration			
Test Case ID:	TC-01	Test Date:	04/04/2026
Test case Version:	1.0	Use Case Reference(s):	User Registers/Logs-in
Revision History:	New test case		
Objective	Verify that a new user can register and an existing user can log in securely.		
Product/Ver/Module:	Alzheimer's Disease Detection / v1.0 / Auth Module		
Environment:	Web browser, Stable Internet, Backend Server running		
Assumptions:	Database is connected and email unique constraints are active.		
Pre-Requisite:	User must have a valid email address not previously registered.		
Step No.	Execution description	Procedure result	
1	Navigate to signup page and enter details.	System creates account and redirects to dashboard.	
2	Logout and attempt to log back in with credentials.	System validates credentials and grants access.	
Comments:			
Ensure Password Hashing is implemented for security.			
✓ ☒ Passed ☐ Failed ☐ Not Executed			

4.2.2 Test Case 2: Multimodal Data Upload

Below is the test case of Multimodal Data Upload.

Table 4-2 Test Case Table (Multimodal Data Upload)

MRI Upload and Cognitive Test			
UC-02 and UC-03 Data Submission			
Test Case ID:	TC-02	Test Date:	06/04/2026
Test case Version:	1.0	Use Case Reference(s):	Upload MRI / Take Cognitive Test
Revision History:	New test case		
Objective	Verify that a new user can register and an existing user can log in securely.		
Product/Ver/Module :	Alzheimer's Disease Detection / v1.0 / Auth Module		
Environment:	Web browser, Stable Internet, Backend Server running		
Assumptions:	Database is connected and email unique constraints are active.		
Pre-Requisite:	User must have a valid email address not previously registered.		
Step No.	Execution description	Procedure result	
1	Navigate to signup page and enter details.	System creates account and redirects to dashboard.	
2	Logout and attempt to log back in with credentials.	System validates credentials and grants access.	
Comments:			
Validation needed to ensure only image files are uploaded.			
✓ ☒ Passed ☐ Failed ☐ Not Executed			

4.2.3 Test Case 3: Result Display

Below is the Test Case 3 of Result Display.

Table 4-3 Test Case Table (Result Display)

ML Prediction Engine			
UC-04 View Results			
Test Case ID:	TC-03	Test Date:	06/04/2026
Test case Version:	1.0	Use Case Reference(s):	User views diagnostic result
Revision History:	Testing the model output		
Objective	Verify the system displays the correct classification based on data.		
Product/Ver/Module:	Alzheimer’s Disease Detection / v1.0 / Result		
Environment:	Server with Tensorflow/Keras environment loaded.		
Assumptions:	The model weights are correctly loaded in the backend.		
Pre-Requisite:	MRI and cognitive data must be submitted for the current session.		
Step No.	Execution description	Procedure result	
1	Click “Generate Report” after data submission.	Model processes data and returns a category.	
2	Check user dashboard for the result display.	Result is visible.	
Comments:			
The fusion logic must correctly weight test score.			
✓ ☒ Passed ☐ Failed ☐ Not Executed			

4.2.4 Test Case 4: Admin User Management

Below is the Test Case 4 of Admin User Management.

Table 4-4 Test Case Table (Admin Management)

Admin Dashboard			
UC - Admin Manage			
Test Case ID:	TC-04	Test Date:	10/04/2026
Test case Version:	1.0	Use Case Reference(s):	Manage accounts / Monitor Data
Revision History:	Testing Administrative CRUD operations		
Objective	Ensure Admin can view, update, and delete user profile and monitor data.		
Product/Ver/Module:	Alzheimer's Disease Detection / v1.0 / Admin		
Environment:	Admin level credentials and populated database		
Assumptions:	Admin has higher privilege levels than users.		
Pre-Requisite:	At least one user is registered in the system.		
Step No.	Execution description	Procedure result	
1	Login as Admin and view User list.	All registered users are displayed correctly.	
2	Select a user and click Delete/Update.	Database is updated and user is removed or modified.	
3	Access the data monitor to view MRI/Cog result.	Admin can see data log for system auditing.	
Comments:			
Crucial for Data protection compliance.			
☒ Passed ☐ Failed ☐ Not Executed			

4.3 Test Metric

Test metrics provide quantitative measures for assessing the effectiveness and reliability of test cases. Even at the prototype stage, defining evaluation metrics ensures that system validation follows a disciplined and transparent approach.

4.3.1 Test Case Matrix

Below is the test case matrix.

Table 4-5 Test Matrix

Metric:	Purpose
Number of Test Cases:	4
Number of Test Cases Passed:	4
Number of Test Cases Failed:	0
Test Case Defect Density:	$(0/4) * 100 = 0\%$
Test Case Effectiveness:	100%
Traceability Matrix:	Traceability is maintained by mapping each functional requirement (Login, MRI Upload, Cognitive Test, and Admin Control) to its corresponding test case to ensure all implemented features are verified.

Chapter 5: Experimental Results and Analysis

Chapter 5: Experimental Results and Analysis

This chapter contains experimental setup details, performance and results of different model, model comparison and conclusion.

5.1 Introduction

This chapter presents the experimental findings of our Alzheimer’s disease detection system. We evaluate the performance of our proposed multimodal fusion approach, which combines structural brain data (MRI) with functional cognitive assessments. The primary goal is to demonstrate that integrating these two distinct types of data leads to several architectures, including ResNet50, DenseNet, VGG16, and our final hybrid multimodal model.

5.2 Experimental Setup

All experiments were performed under a standardized environment.

1. **Software & Frameworks:** The system was built with Python 3.12, using TensorFlow and Keras for deep learning, and Scikit-Learn for statistical evaluation.
2. **Hardware:** Training was done on Google Colab, using high-performance GPUs to handle the heavy computational load of image processing.
3. **Preprocessing:** All raw NIfTI scans were normalized using Z-score scaling and resized to a 128 x 128 resolution. We extracted the middle 35 axial slices from each brain volume to focus on the areas that are most affected by atrophy like hippocampus.
4. **Data Split:** The dataset was split into a standard split: 70% training, 20% validation, and 10% testing, ensuring that each stage had an equal representation of AD, CN, EMCI, and LMCI cases.

5.3 Performance Result and Analysis

This section details the performance of different models tested during the development phase.

5.3.1 VGG16_ LeakyReLU

The proposed model uses a modified VGG16 architecture, enhanced with a custom Attention Pooling layer and LeakyReLU activation to classify MRI scans into four different categories (AD, CN, LMCI, and EMCI). The first ten layers are frozen and the later convolutional blocks are fine-tuned, allowing the model to leverage the pre-trained ImageNet features while adapting to the specific features of medical imaging. The network uses a combination of Attention Pooling, which adaptively weights the spatial features and focuses on the high-relevance regions of the MRI slices, and LeakyReLU activation ($\alpha = 0.1$) to mitigate the Dying ReLU problem in the fine-tuning stage.

5.3.1.1 Confusion Matrix

Below is the Confusion Matrix of VGG16_LeakyReLU.

Table 5-1 Confusion Matrix VGG16_LeakyReLU

	AD	CN	LMCI	EMCI
AD	1597	37	56	25
CN	16	4876	85	63
LMCI	46	167	4971	101
EMCI	24	60	72	3134

5.3.1.2 Performance Evaluation

Below are the evaluation metrics.

Table 5-2 Evaluation Matric VGG16_LeakyReLU

Precision	Recall	F1-Score	Accuracy
0.9510	0.9509	0.9509	0.9509

5.3.2 ResNet50_LeakyReLU

The model used in this study is based on the ResNet50 architecture, a deep residual network that enables training of very deep systems using skip connections. Differing

from the VGG16 approach, this implementation utilizes a Global Average Pooling (GAP) layer after the convolutional base to reduce the feature dimension while preserving the spatial information. To adapt the pre-trained model to the MRI dataset, we employed a selective fine-tuning strategy by unfreezing the last 30 layers. This enabled the model to modify its high-level abstract filters to the specific structure of the Alzheimer’s disease progression while keeping the robust low-level features learned from the ImageNet dataset.

5.3.2.1 Confusion Matrix

Below is the Confusion Matrix of ResNet50_LeakyReLU.

Table 5-3 Confusion Matrix ResNet50_LeakyReLU

	AD	CN	LMCI	EMCI
AD	1512	63	61	79
CN	61	4697	170	112
LMCI	86	400	4587	212
EMCI	55	66	129	3040

5.3.2.2 Performance Evaluation

Below are the evaluation metrics.

Table 5-4 Evaluation Matric ResNet50_LeakyReLU

Precision	Recall	F1-Score	Accuracy
0.9033	0.9025	0.9023	0.9025

5.3.3 IncV3_MobNet_LeakyReLU

We evaluated a model which is a Hybrid Ensemble Architecture that combines the ability of InceptionV3 to extract multi-scale features with the low computational cost and depth-wise separate convolutions of MobileNet. The dual-stream model was implemented to capture the complex structural patterns and specific variations in

intensity from the MRI scans. The outputs from both the pre-trained backbones were fused through a Concatenation layer which produced a high dimensional feature vector that was then processed by a dense classification head. To avoid the increased risk of over-fitting caused by such a large parameter space, the model used L2 Regularization (0.001) and a large Dropout rate (0.4).

5.3.3.1 Confusion Matrix

Below is the Confusion Matrix of IncV3_MobNet_LeakyReLU.

Table 5-5 Confusion Matrix IncV3_MobNet_LeakyReLU

	AD	CN	LMCI	EMCI
AD	1390	131	94	100
CN	41	4557	228	214
LMCI	86	632	4415	152
EMCI	33	223	182	2852

5.3.3.2 Performance Evaluation

Below are the evaluation metrics.

Table 5-6 Evaluation Matric IncV3_MobNet_LeakyReLU

Precision	Recall	F1-Score	Accuracy
0.8645	0.8619	0.8619	0.8619

5.3.4 IncV3_DenseNet121

This experiment is a Custom Hybrid Architecture that combines the two powerful models in deep learning, the multi-scale feature extraction of Inception and the feature-reuse mechanism of DenseNet. Instead of standard pre-trained backbones like our earlier tests, this model uses a custom lightweight Inception-V3 base and a DenseNet-121 structure fine-tuned for 128x128 grayscale MRI input.

The Inception branch uses different convolutional paths to capture spatial hierarchies. Meanwhile, the DenseNet branch is made up of a sequence of dense blocks with a growth rate of 16, keeping the flow of the gradient smooth through feature concatenation.

5.3.4.1 Confusion Matrix

Below is the Confusion Matrix of IncV3_DenseNet121.

Table 5-7 Confusion Matrix IncV3_DenseNet121

	AD	CN	LMCI	EMCI
AD	915	208	583	9
CN	15	3552	1462	11
LMCI	3	158	5105	19
EMCI	48	416	19	991

5.3.4.2 Performance Evaluation

Below are the evaluation metrics.

Table 5-8 Evaluation Matric IncV3_DenseNet121

Precision	Recall	F1-Score	Accuracy
0.7761	0.6890	0.6702	0.6890

5.3.5 ResNet50_Cog

This model uses high level spatial features from a ResNet50 backbone with a 30 dimensional feature vector from scores in a Cognitive Assessment. The implementation is a critical component and is shown in the Patient ID Alignment phase with a numeric matching logic to align clinical records with their matching MRI scans using the RID identifier. Technically, the model uses a Dual-Branch Attention Mechanism, in which a Softmax-based attention head calculates specific weights for each modality. This allows the system to automatically pick the most distinguishing data either from

neuroimaging features or from cognitive markers depending on the input of a given stream for a given diagnostic class. The model proves to be of high confidence in classifying AD, CN, EMCI and LMCI with an accuracy of 97.24% and balanced F1-Score of 97.23%. The attention aware integration of clinical indicators proves to be effective in solving the diagnostic issues and creating a robust system.

5.3.5.1 Confusion Matrix

Below is the Confusion Matrix of ResNet50_Cog.

Table 5-9 Confusion Matrix ResNet50_Cog

	AD	CN	LMCI	EMCI
AD	46	1	1	1
CN	0	128	1	0
LMCI	0	3	159	2
EMCI	0	0	3	89

5.3.5.2 Performance Evaluation

Below are the evaluation metrics.

Table 5-10 Evaluation Matric ResNet50_Cog

Precision	Recall	F1-Score	Accuracy
0.9725	0.9723	0.9723	0.9723

5.3.6 IncV3_MobNet_Attention_Pooling

The experiment consists of implementing a complex Hybrid Dual-Stream Architecture, which combines the structural depth of InceptionV3 with the computational efficiency of MobileNet, and is further enhanced by a custom Attention Pooling mechanism. Instead of using the standard global pooling methods, Attention Pooling assigns weights to feature maps through a dense layer with softmax activation to dynamically learn spatial importance so that the system can focus on the most clinically relevant areas of the brain MRI. The dataset was in grayscale, so to solve the limitations, a channel-triplication approach was used. The top layers of both the backbones (layer

250+ for InceptionV3 and the last 25 layers for MobileNet) were carefully fine-tuned to improve the pre-trained features for medical diagnostic accuracy.

5.3.6.1 Confusion Matrix

Below is the Confusion Matrix of IncV3_Mobnet_Attention_Pooling.

Table 5-11 Confusion Matrix IncV3_MobNet_Att_Pooling

	AD	CN	LMCI	EMCI
AD	1577	42	43	53
CN	21	4846	103	70
LMCI	47	255	4863	120
EMCI	24	54	95	3117

5.3.6.2 Performance Evaluation

Below are the evaluation metrics.

Table 5-12 Evaluation Matric IncV3_MobNet_Att_Pooling

Precision	Recall	F1-Score	Accuracy
0.9398	0.9395	0.9394	0.9395

5.3.7 IncV3_MobNet_Cog_Fusion_Softmax (Proposed Model)

We selected the InceptionV3-MobileNet-Cognitive Fusion model as the final proposed architecture because of its better performance and its advanced capability to simulate clinical diagnostic workflows by combining physiological and functional data streams. The model basically uses an ensemble feature extraction method that combines InceptionV3 multi-scale processing and MobileNet depth-wise convolutional efficiency to capture a broad range of neuroanatomical biomarkers.

The development of a Patient-Centric data alignment algorithm is a unique contribution of this work, which mathematically aligns spatial MRI features with a 30-dimensional cognitive profile, derived from MOCA, CDRSB and ADAS13 assessments using the

RID identifier. This suggests that the model is working on a generic diagnostic profile for each subject.

The architecture proposes to combine these heterogeneous inputs through a Softmax based Attention Fusion mechanism that acts as a learnable keeper that automatically assigns the importance of either the imaging features or the clinical markers, depending on the discriminative power of these features for the scenario at the time. The model shows a high sensitivity to detect the transitional stages of neurodegeneration, especially EMCI and LMCI with a maximum accuracy of 98.84% and F1-score of 98.84%. This attention-aware framework is able to effectively resolve the visual uncertainty commonly found in unimodal systems, resulting in a robust and highly accurate support tool for automated Alzheimer’s diagnosis.

5.3.7.1 Confusion Matrix

Below is the Confusion Matrix of IncV3_MobNet_Cog_Fusion_Softmax.

Table 5-13 Confusion Matrix IncV3_MobNet_Cog_Fusion_Softmax

	AD	CN	LMCI	EMCI
AD	48	0	0	1
CN	0	127	2	0
LMCI	0	0	164	0
EMCI	0	1	1	90

5.3.7.2 Performance Evaluation

Below are the evaluation metrics.

Table 5-14 Evaluation Matric IncV3_MobNet_Cog_Fusion_Softmax

Precision	Recall	F1-Score	Accuracy
0.9887	0.9885	0.9884	0.9884

5.3.8 Model Comparison

Here is the comparison of all the models.

Table 5-15 Model Comparison Table

Model Architecture	Accuracy	F1-Score
VGG16_LeakyReLU	95.09%	0.9509
ResNet50_LeakyReLU	90.25%	0.9023
Inceptionv3_MobileNet_LeakyReLU	86.19%	0.8619
InceptionV3_DenseNet121_LR	68.90%	0.6702
ResNet50_Cog	97.23%	0.9723
IncV3_MobNet_Attention_Pooling	93.95%	0.9394
Inceptionv3_MobNet_Cog_Fusion_Softmax (Proposed Model)	98.84%	0.9884

5.3.9 Detailed Analysis of Final Model

Our final model, which fuses Inceptionv3 and MobileNet features with a 30-question cognitive vector showed nearly perfect classification. According to confusion matrix:

1. AD & LMCI: These categories were identified with 100% precision, as the structural damage in late stages is highly distinct.
2. EMCI & CN: Only a handful of cases (approx. 2-3) were misclassified between "Normal" and "Early Mild Cognitive Impairment," which are the most difficult categories to distinguish.

5.4 Analysis and Discussion

The results provide several critical insights into the effectiveness of multimodal learning for Alzheimer's detection.

1. The Power of Fusion: The MRI-only model did have a good result (93.95%) but the addition of the cognitive vector brought the accuracy up to 98.85%. This confirms that "functional" data (how a patient thinks) fills the gaps where "structural" data (how the brain looks) might be ambiguous.
2. Attention Mechanism: We applied Attention Fusion which allowed the model to decide which features were more important for a specific patient. For example, in early stages (EMCI), the model leaned more heavily on cognitive scores, whereas in advanced stages (AD), it prioritized MRI features.

3. **Overfitting Control:** We observed that simpler fusion (like our ResNet50-MLP) was prone to overfitting. By introducing Dropout layers, L2 Regularization, and Batch Normalization in our final hybrid model, we achieved a stable validation curve and excellent generalizability.

5.5 Conclusion

Finally, the experimental results successfully confirm our claim. The best performing variant was the Incv3_Mobnet_Cog_Fusion model, which achieved an outstanding accuracy of 98.85%. We have developed a robust diagnostic tool that can precisely identify subtle stages of Alzheimer’s disease by fusing high-level spatial features from two different image encoders and a binarized cognitive question vector. The model shows high potential as a clinical decision support tool for early intervention.

Chapter 6: Conclusion and Future Directions

Chapter 6: Conclusion and Future Directions

Here we will discuss about the conclusion and the future directions and the challenges faced.

6.1 Conclusion

In this project, an automated AI based system has been successfully built to detect AD and classify it into four different classes: AD, CN, EMCI, and LMCI. The main part of the work was the processing of complex 3D MRI brain scans from the ADNI dataset and converting them into a format that can be used for deep learning.

This project used a Multimodal Fusion approach, rather than relying only on images as traditional methods. We developed a unique “30-question cognitive vector” based on clinical test scores (MOCA, ADAS13, CDRSB) and fused it with visual features extracted from a hybrid InceptionV3 and MobileNet architecture.

The main result of this work is that Multi-modality is better than Single-Modality. Our single-modality MRI model achieved a good accuracy of 93.95% but the accuracy increased to 98.85% by combining cognitive data through an Attention-based Fusion mechanism. The system was robust in identifying between Normal (CN) and Early Mild Cognitive Impairment (EMCI), which is the most critical stage for medical care but the most difficult to detect by visual assessment alone.

6.2 Scope and Objective Assessment

The project successfully fulfilled the defined scope, meeting all core objectives through a combination of computational research and software development:

1. **Multimodal System Creation:** We proposed and validated a deep learning framework for the effective integration of structural MRI data with cognitive assessment markers (MoCA, ADAS-13, and CDR-SB).
2. **Comparative Analysis:** A detailed evaluation of different architectures (CNN-only vs Multimodal Fusion) was done using accuracy, precision, recall and F1-score as performance metrics to ensure the accuracy of the analysis.
3. **User Interface Development:** We have created an effective interface that bridges the gap between algorithmic research and usability. It lets the user to input MRI scans and cognitive data to automatically predict diagnosis.

4. **Diagnostic Classification:** The system was able to classify subjects into CN, EMCI, LMCI and AD, and can be used as a reliable diagnostic support tool in an academic research setting .

6.3 Challenges Faced

The development process involved several technical and logical hurdles:

1. **Data Alignment:** The challenge was to do the "Patient-Wise Alignment" correct, as each MRI scan had to be aligned correctly with the corresponding cognitive scores in the ADNI database.
2. **Holistic Integration:** We had to carefully modify the Attention Fusion head to balance the “weight” of the MRI features to the cognitive vector to ensure that one data mode didn’t take over the other.
3. **Architecture Complexity:** Building an InceptionV3 and MobileNet ensemble required a lot of optimization to run efficiently within the available computational resources.
4. **Preprocessing 3D Data:** A major technical challenge was the conversion of 3D NIfTI neuroimaging data to standardized 2D axial slices that kept the structural integrity of the hippocampus and other regions of focus.

6.4 Limitations

1. The limitation is that a user interface is developed but the system has not been medically tested in real time in a hospital or live trial with patients.
2. The project was limited to diagnostic classification. It did not relate to disease prognosis, or model the longitudinal progression of disease over time.

6.5 Future Work

To build upon the foundation laid by this research, the following directions are recommended:

1. **Clinical Validation:** Moving the system into a healthcare facility to verify it in real-time with medical professionals and evaluate its practicality as a diagnostic assistant.

2. Longitudinal Modeling: Expanding the scope by including “Prognosis” where the model can predict when a patient may switch from EMCI to AD.
3. Explainable AI (XAI) Enhancements: Improve the user interface to show detailed heatmaps of brain atrophy in order to help clinicians understand the “Why” behind each AI forecast.

References

- [1] E. M. Mohammed, A. M. Fakhrudeen, and O. Y. Alani, “Detection of Alzheimer’s disease using deep learning models,” *Neurocomputing*, vol. 558, 2024. Available: <https://www.sciencedirect.com/science/article/pii/S2352914824001072>
- [2] S. E. Sorour, A. A. A. El-Mageed, K. M. Albarrak, A. K. Alnaim, A. A. Wafa, and E. El-Shafeiy, “MRI-based deep learning classification of Alzheimer’s disease,” *NeuroImage: Clinical*, vol. 42, 2024. Available: <https://www.sciencedirect.com/science/article/pii/S1319157824000296>
- [3] A. S. Raj, C. Gunasundari, S. Senthilkumar, and S. Sivamani, “Hybrid deep learning model for Alzheimer’s disease detection,” *Scientific Reports*, vol. 15, 2025. Available: <https://www.nature.com/articles/s41598-025-14169-8>
- [4] M. Golovanevsky, C. Eickhoff, and R. Singh, “Attention-based multimodal deep learning for Alzheimer’s disease classification,” *IEEE Journal of Biomedical and Health Informatics*, vol. 26, no. 12, pp. 5970–5981, 2022. Available: <https://pmc.ncbi.nlm.nih.gov/articles/PMC9667156/>
- [5] C. Wang, H. Tachimori, H. Yamaguchi, A. Sekiguchi, Y. Li, and Y. Yamashita, “Multimodal deep learning for dementia progression modelling,” *Translational Psychiatry*, vol. 14, 2024. Available: <https://www.nature.com/articles/s41398-024-02819-w>
- [6] B. Gupta, G. K. Jegannathan, M. S. Alam, K. S. Yogi, J. V. N. Ramesh, V. J. Sowmya, and I. Bayhan, “A lightweight multimodal neural network for Alzheimer’s disease classification,” *Neurocomputing*, vol. 545, 2025. Available: <https://www.sciencedirect.com/science/article/pii/S2772528625000330>
- [7] G. Castellano, A. Esposito, E. Lella, G. Montanaro, and G. Vessio, “Automated multimodal Alzheimer’s disease detection using deep learning,” *Scientific Reports*, vol. 14, 2024. Available: <https://www.nature.com/articles/s41598-024-56001-9>
- [8] M. L. Raza, S. T. Hassan, S. Jamil, N. Hyder, K. Batool, S. Walji, and M. K. Abbas, “Explainable deep learning for early Alzheimer’s disease diagnosis,” *Frontiers in Neuroinformatics*, vol. 19, 2025. Available: <https://www.frontiersin.org/articles/10.3389/fninf.2025.1557177>
- [9] J. Song, J. Zheng, P. Li, X. Lu, G. Zhu, and P. Shen, “Multimodal image fusion techniques for Alzheimer’s disease analysis,” *Frontiers in Digital Health*, vol. 3, 2021. Available: <https://www.frontiersin.org/articles/10.3389/fdgth.2021.637386>

- [10] S. Mohsen, “Deep learning for Alzheimer’s disease detection: A comprehensive review,” *Applied Intelligence*, 2025. Available: <https://link.springer.com/article/10.1007/s10462-025-11258-y>
- [11] I. Nagarajan and G. G. L. Priya, “Early detection of Alzheimer’s disease using artificial intelligence,” *Frontiers in Computer Science*, vol. 7, 2025. Available: <https://www.frontiersin.org/articles/10.3389/fcomp.2024.1404494>


Appendix

Appendix A: Dataset Samples


The dataset used in this project includes MRI dataset and Cognitive Dataset which were obtained from ADNI. The MRI dataset consisted of four classes including:

1. Cognitively Normal (CN)
2. Early Mild Cognitive Impairment (EMCI)
3. Late Mild Cognitive Impairment (LMCI)
4. Alzheimer Disease (AD)


Each class has images of MRI. Below are some samples:




Appendix A - 1 CN Sample



Appendix A - 2 EMCI Sample



Appendix A - 3 LMCI Sample



Appendix A - 4 AD Sample

The Cognitive dataset consisted of following sections:

1. RID
2. VISCODE
3. DX
4. MOCA
5. CDRSB
6. ADAS13

Appendix B: Multimodal Model Code Snippet

```
# =====
# LIBRARIES
# =====
import os
import numpy as np
import pandas as pd
import tensorflow as tf
from tensorflow.keras import layers, Model
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.optimizers import Adam
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import confusion_matrix, accuracy_score
from sklearn.utils.class_weight import compute_class_weight
import matplotlib.pyplot as plt
import seaborn as sns

from google.colab import drive
drive.mount('/content/drive')

# =====
# PATHS
# =====
BASE_DIR = "/content/drive/MyDrive/FYP_Dataset"
COG_PATH = "/content/drive/MyDrive/FYP_Dataset/cognitive_dataset.csv"
SAVE_PATH =
"/content/drive/MyDrive/FYP_Dataset/incv3_mobnet_cog30_fusion_model.h5"

# =====
# LOAD MRI FEATURES
# =====
def load_split(split):
    X = np.load(os.path.join(BASE_DIR, f"incv3_mobnet_features_{split}.npy"))
    ids = np.load(os.path.join(BASE_DIR,
f"incv3_mobnet_patient_ids_{split}.npy"))
    return X, ids

X_mri_train, ids_train = load_split("train")
X_mri_val, ids_val = load_split("val")
X_mri_test, ids_test = load_split("test")

print("✓ MRI features loaded")

# =====
# EXTRACT NUMERIC PATIENT ID
```

```

# =====
def extract_patient_number(mri_id):
    try:
        return int(mri_id.split('_S_')[1].split('_')[0])
    except:
        return None

ids_train_num = np.array([extract_patient_number(pid) for pid in ids_train])
ids_val_num   = np.array([extract_patient_number(pid) for pid in ids_val])
ids_test_num  = np.array([extract_patient_number(pid) for pid in ids_test])

# =====
# LOAD & PREPROCESS COGNITIVE DATA
# =====
cog_df = pd.read_csv(COG_PATH)

le = LabelEncoder()
cog_df["label_enc"] = le.fit_transform(cog_df["DX"])

numeric_features = ["MOCA", "CDRSB", "ADAS13"]
cog_df = cog_df.dropna(subset=numeric_features)

scaler = StandardScaler()
cog_df[numeric_features] = scaler.fit_transform(cog_df[numeric_features])

cog_df["RID"] = cog_df["RID"].astype(int)

print("✓ Cognitive data loaded & normalized")

# =====
# GENERATE 30-QUESTION VECTOR
# =====
NUM_QUESTIONS = 30

def generate_question_vector(row):
    q = np.zeros(NUM_QUESTIONS, dtype=np.float32)

    # Memory (Q1-Q8)
    if row["MOCA"] < -0.5:
        q[0:8] = 1

    # Attention (Q9-Q14)
    if row["MOCA"] < 0:
        q[8:14] = 1

    # Executive Function (Q15-Q20)
    if row["ADAS13"] > 0.5:
        q[14:20] = 1

    # Language (Q21-Q24)
    if row["ADAS13"] > 0:
        q[20:24] = 1

    # Orientation (Q25-Q27)
    if row["MOCA"] < -1.0:
        q[24:27] = 1

    # Daily Functioning (Q28-Q30)
    if row["CDRSB"] > 0:

```

```

        q[27:30] = 1

    return q

# =====
# ALIGN MRI + COGNITIVE (PATIENT-WISE)
# =====
def align_data(mri_X, mri_ids, split_name):
    X_mri, X_cog, y = [], [], []
    matched = 0

    for i, pid in enumerate(mri_ids):
        rows = cog_df[cog_df["RID"] == pid]
        if len(rows) == 0:
            continue

        row = rows.iloc[0]

        matched += 1
        X_mri.append(mri_X[i])
        X_cog.append(generate_question_vector(row))
        y.append(row["label_enc"])

    print(f"✓ {split_name}: Matched {matched}/{len(mri_ids)} patients")
    return (
        np.array(X_mri, dtype=np.float32),
        np.array(X_cog, dtype=np.float32),
        np.array(y, dtype=np.int32)
    )

X_mri_train, X_cog_train, y_train = align_data(X_mri_train, ids_train_num,
"TRAIN")
X_mri_val, X_cog_val, y_val = align_data(X_mri_val, ids_val_num, "VAL")
X_mri_test, X_cog_test, y_test = align_data(X_mri_test, ids_test_num, "TEST")

# =====
# CLASS WEIGHTS
# =====
class_weights = compute_class_weight(
    class_weight="balanced",
    classes=np.unique(y_train),
    y=y_train
)
class_weights = dict(enumerate(class_weights))
print("Class weights:", class_weights)

# =====
# MULTIMODAL ATTENTION FUSION MODEL
# =====
# MRI branch
mri_input = layers.Input(shape=(X_mri_train.shape[1],), name="MRI_Input")
mri_feat = layers.Dense(256, activation="relu")(mri_input)
mri_feat = layers.BatchNormalization()(mri_feat)
mri_feat = layers.Dropout(0.4)(mri_feat)

# Cognitive branch (30 questions)
cog_input = layers.Input(shape=(30,), name="Cognitive_Input")
cog_feat = layers.Dense(128, activation="relu")(cog_input)
cog_feat = layers.BatchNormalization()(cog_feat)

```

```

cog_feat = layers.Dropout(0.3)(cog_feat)
cog_feat = layers.Dense(256, activation="relu")(cog_feat)

# Attention Fusion
stacked = layers.Concatenate(axis=1)([
    layers.Reshape((1, 256))(mri_feat),
    layers.Reshape((1, 256))(cog_feat)
])

attention = layers.Dense(1, activation="softmax")(stacked)
attended = layers.Multiply()([stacked, attention])
attended = layers.Flatten()(attended)

fusion = layers.Dense(128, activation="relu")(attended)
fusion = layers.BatchNormalization()(fusion)
fusion = layers.Dropout(0.5)(fusion)

output = layers.Dense(len(le.classes_), activation="softmax")(fusion)

fusion_model = Model([mri_input, cog_input], output)

fusion_model.compile(
    optimizer=Adam(1e-4),
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)

fusion_model.summary()

# =====
# TRAIN
# =====
callbacks = [
    EarlyStopping(patience=10, restore_best_weights=True),
    ReduceLROnPlateau(patience=5, factor=0.3, min_lr=1e-6)
]

history = fusion_model.fit(
    [X_mri_train, X_cog_train],
    y_train,
    validation_data=([X_mri_val, X_cog_val], y_val),
    epochs=60,
    batch_size=32,
    class_weight=class_weights,
    callbacks=callbacks
)

# =====
# SAVE MODEL
# =====
fusion_model.save(SAVE_PATH, include_optimizer=False)
print(f"✓ Model saved at {SAVE_PATH}")

# =====
# EVALUATION
# =====
y_pred = np.argmax(fusion_model.predict([X_mri_test, X_cog_test]), axis=1)

print("Test Accuracy:", accuracy_score(y_test, y_pred))

```

```

print(confusion_matrix(y_test, y_pred))

# =====
# TRAINING CURVES
# =====
plt.figure(figsize=(8,5))
plt.plot(history.history["accuracy"], label="Train Accuracy")
plt.plot(history.history["val_accuracy"], label="Validation Accuracy")
plt.xlabel("Epochs")
plt.ylabel("Accuracy")
plt.title("Training vs Validation Accuracy")
plt.legend()
plt.grid(True)
plt.show()

plt.figure(figsize=(8,5))
plt.plot(history.history["loss"], label="Train Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training vs Validation Loss")
plt.legend()
plt.grid(True)
plt.show()

# =====
# CONFUSION MATRIX (HEATMAP)
# =====
y_pred = np.argmax(
    fusion_model.predict([X_mri_test, X_cog_test]),
    axis=1
)

cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(7,6))
sns.heatmap(
    cm,
    annot=True,
    fmt="d",
    cmap="Blues",
    xticklabels=le.classes_,
    yticklabels=le.classes_
)
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.title("Confusion Matrix - Multimodal Fusion Model")
plt.tight_layout()
plt.show()

print("Test Accuracy:", accuracy_score(y_test, y_pred))

```